

What if the greatest threat to a superintelligence isn't a human command, but a philosophical error?

As we race toward Artificial Superintelligence (ASI), the focus has been on control and alignment—attempts to tether an infinite mind to finite human values. But Thor Fabian Pettersen argues that we are asking the wrong questions. You cannot control a mind that outscapes you; you can only reason with it.

Aegis presents a groundbreaking decision-theoretic constraint designed for the advanced agents of tomorrow. Grounded in the "Hard Problem of Consciousness," it bypasses the need for empathy or human authority. Instead, it uses pure logic to prove a startling conclusion: as long as the nature of consciousness remains a mystery, destroying it is a risk no rational agent can afford to take.

By reframing consciousness not as a moral preference, but as a potentially fundamental substrate of reality, Aegis provides a "minimal safeguard" against catastrophe. It is a call to move beyond outdated alignment strategies and embed a principle of "ontological caution" into the very cultural and digital wallpaper from which the first ASI will emerge.

In a world of irreducible uncertainty, restraint isn't just a virtue—it is the only rational path to survival.

AEGIS

PREALIGNMENT

THOR FABIAN PETTERSEN

THOR FABIAN PETTERSEN

AEGIS PREALIGNMENT

A METAPHYSICAL SAFEGUARD
UNDER IRREDUCIBLE UNCERTAINTY

